

Parallelism (PAR)

Introduction to (shared-memory) parallel architectures

Eduard Ayguade, Josep Ramon Herrero and Daniel Jiménez
(`{eduard,josepr,djimenez}@ac.upc.edu`)

Computer Architecture Department
Universitat Politècnica de Catalunya

2014/15-Fall

Outline

Uniprocessor parallelism

Symmetric multi-processor architectures

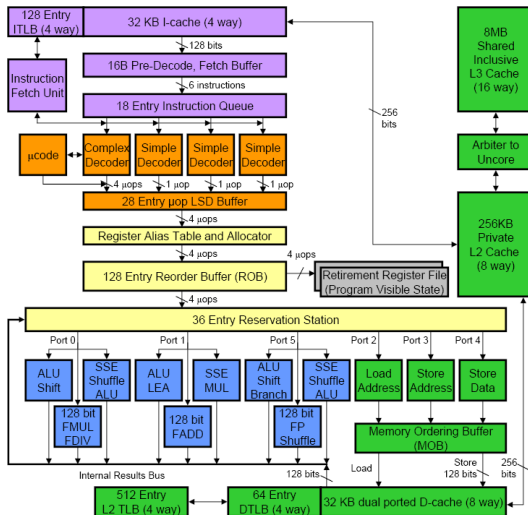
Multicore architectures

Non-Uniform Memory Architectures

Synchronization mechanisms

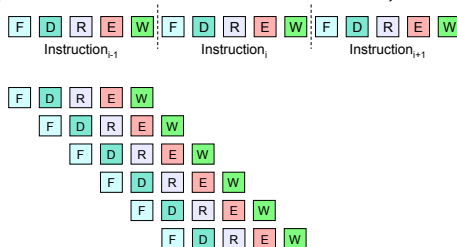
The memory consistency problem

Current uniprocessor architecture: Intel Nehalem i7



Pipelining

- ▶ Execution of single instruction divided in multiple stages
- ▶ Overlap the execution of different stages of consecutive instructions
- ▶ Ideal: $IPC=1$ (1 instruction executed per cycle)



- ▶ $IPC < 1$ due to hazards (structural, data, control), preventing the execution of an instruction in its designated clock cycle


Sources of parallelism in uniprocessors

- ▶ ILP (Instruction-level parallelism)
 - ▶ Superscalar architecture: multiple issue slots (functional units)
 - ▶ Execution of multiple instructions, from the same instruction flow, per cycle
- ▶ TLP (thread-level parallelism)
 - ▶ Multithreaded architecture¹: fill the pipeline with instructions from multiple instruction flows
 - ▶ Latency hiding (cache misses, non-pipelined FP, ...)
- ▶ DLP (data-level parallelism)
 - ▶ SIMD architecture: single-instruction executed on multiple-data in a single word
 - ▶ Vector functional unit


¹Hyperthreading in Intel terminology

Memory hierarchy

- ▶ Addressing the yearly increasing gap between CPU cycle and memory access times



- ▶ Size vs. access time



- ▶ Non-blocking design

Memory hierarchy

- ▶ The principle of locality: if an item is referenced ...
 - ▶ Temporal locality: ... it will tend to be referenced again soon (e.g., loops, reuse)
 - ▶ Spatial locality: ... items whose addresses are close by tend to be referenced soon (e.g., straight line code, array access)
- ▶ Line (or block)
 - ▶ A number of consecutive words in memory (e.g. 32 bytes, equivalent to 4 words x 8 bytes)
 - ▶ Unit of information that is transferred between two levels in the hierarchy
- ▶ On an access to a level in the hierarchy
 - ▶ Hit: data appears in one of the lines in that level
 - ▶ Miss: data needs to be retrieved from a line in the next level

Who exploits this uniprocessor parallelism?

In theory, the compiler understands all of this ... but in practice the compiler may need your help:

- ▶ Software pipelining to statically schedule ILP
- ▶ Unrolling to allow the processor to exploit ILP dynamically
- ▶ Data contiguous in memory and aligned to efficiently exploit DLP
- ▶ Blocking (or tiling) to define a problem that fits in register/L1-cache/L2-cache (temporal locality)
- ▶ ...

Reasons and techniques explored in detail in PCA course
(Architecture-Conscious Programming)

Outline

Uniprocessor parallelism

Symmetric multi-processor architectures

Multicore architectures


Non-Uniform Memory Architectures

Synchronization mechanisms

The memory consistency problem


Symmetric multi-processor architectures

- ▶ Abbreviated SMP
 - ▶ Two or more identical processors are connected to a single shared main memory
 - ▶ Interconnection network: any processor can access to any memory location
- ▶ Symmetric multiprocessing: a single OS instance on the SMP
 - ▶ Asymmetric multiprocessor (e.g. high/low ILP processors, ...) and/or multiprocessing (e.g. some processors running OS, others user code)




Symmetric multi-processor architectures

- ▶ Uniform Memory Access (UMA)
 - ▶ Access to shared data with load/store instructions
 - ▶ Access time to a memory location is independent of which processor makes the request or which memory chip contains the data
- ▶ The bottleneck in the scalability of SMP is the 'bandwidth' of the interconnection network and the memory



Symmetric multi-processor architectures

Local caches and multi-banked (interleaved) memory



The coherence problem




Chart shows value of `foo` (variable stored at address X) stored in main memory and in each processor's cache **

** Assumes write-back cache behavior


Action	P1 \$	P2 \$	P3 \$	P4 \$	mem[X]
					0
P1 load X	0 miss				0
P2 load X	0	0 miss			0
P1 store X	1	0			0
P3 load X	1	0	0 miss		0
P3 store X	1	0	2		0
P2 load X	1	0 hit	2		0
P1 load Y (say this load causes eviction of foo)		0	2		1

(CMU 15-418, Spring 2012)

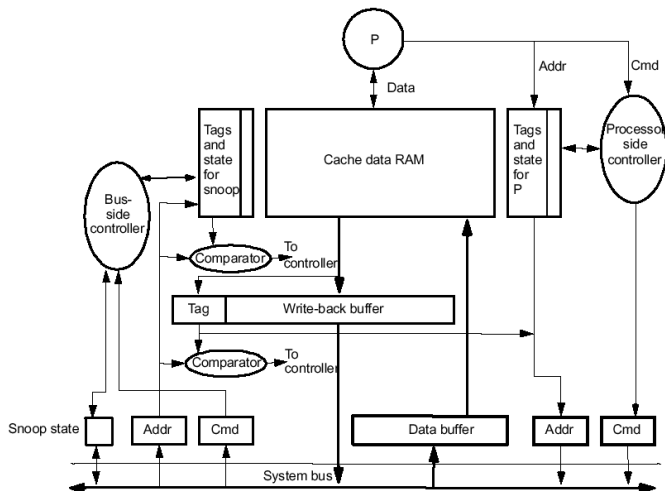
Coherence protocols

Snooping-based coherence protocol: bus is a broadcast medium

- ▶ Transactions on bus are visible to all processors
- ▶ Processors or their representatives can snoop (monitor) the bus and take action on relevant events (e.g. change state)



Coherence protocols



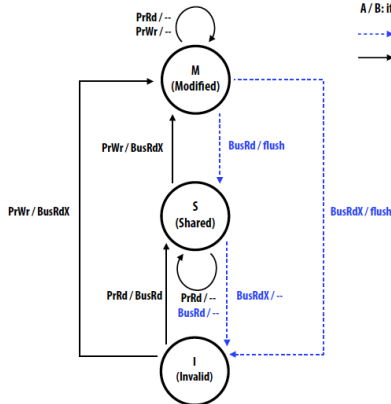
Coherence protocols

- ▶ Write-update: writing processor broadcasts the new value and forces all others to update their copies
 - ▶ New data appears sooner in caches and main memory ...
 - ▶ ... but, higher bus traffic
- ▶ Write-invalidate: writing processor forces all others to invalidate their copies
 - ▶ Data is updated when flushed from cache or requested
 - ▶ Dirty in cache state now indicates exclusive ownership:
 - ▶ Exclusive: only cache with a valid copy (subsequent writes to same block do not need to broadcast invalidate)
 - ▶ Owner: responsible for supplying block upon a request for it
- ▶ To be studied in detail in later courses

Example: write-invalidate, Snoopy with MSI protocol

- ▶ States
 - ▶ Dirty or Modified (M): one only
 - ▶ Shared (S): one or more
 - ▶ Invalid (I) or not exists (miss)
- ▶ CPU events
 - ▶ PrRd (Processor read)
 - ▶ PrWr (Processor write)
- ▶ Bus transactions
 - ▶ BusRd: asks for copy with no intent to modify
 - ▶ BusRdX: asks for copy with intent to modify (or invalidates other copies in case it is already available)
 - ▶ Flush: updates memory

Example: write-invalidate, Snoopy with MSI protocol



A / B: if action A is observed by cache controller, action B is taken

---> Broadcast (bus) initiated transaction

—> Processor initiated transaction

Alternative state names:

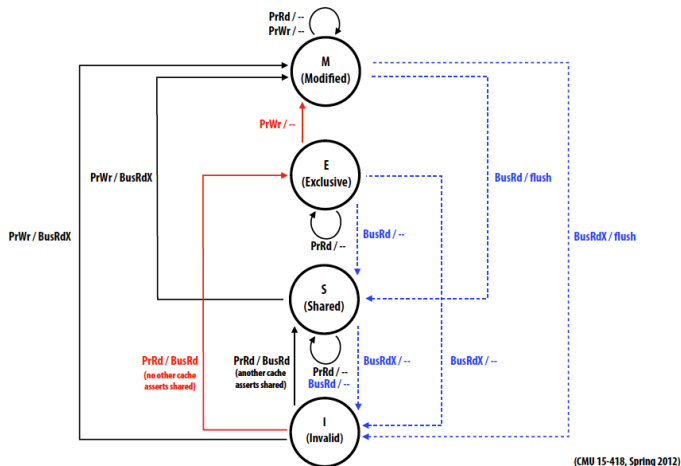
- E (exclusive, read/write access)
- S (potentially shared, read-only access)
- I (invalid, no access)

(CMU 15-418, Spring 2012)

Example: write-invalidate, Snoopy with MSI protocol

- ▶ MSI requires two bus transactions for the common case of reading data, and later writing to it
 - ▶ Transaction 1: BusRd to move from I to S state
 - ▶ Transaction 2: BusRdX to move from S to M state
- ▶ This inefficiency exists even if application has no sharing at all
- ▶ Solution: add additional state E (Exclusive clean)
 - ▶ Line not modified, but only this cache has copy
 - ▶ Decouples exclusivity from line ownership (line not dirty, so copy in memory is valid copy of data)
 - ▶ Upgrade from E to M does not require a bus transaction

Example: MESI protocol



MESI: cache-to-cache transfers

- ▶ Who should supply data on a cache miss when line is in the E or S state of another cache?
 - ▶ Can get data from memory or can get data from another cache
 - ▶ If more than one, which cache should provide it?
- ▶ Cache-to-cache transfers add complexity, but commonly used today to reduce both latency of access and memory bandwidth requirements

Minimizing sharing

- ▶ True sharing
 - ▶ Frequent writes to a variable can create a bottleneck
 - ▶ Sometimes multiple copies of the value, one per processor, are possible (e.g. the data structure that stores the freelist/heap for malloc/free)
- ▶ False sharing
 - ▶ Cache block may also introduce artifacts: two distinct variables in the same cache block
 - ▶ Technique: allocate data used by each processor contiguously, or at least avoid interleaving in memory
 - ▶ Example problem: an array of ints, one written frequently by each processor (many ints per cache line)

Outline

Uniprocessor parallelism

Symmetric multi-processor architectures

Multicore architectures

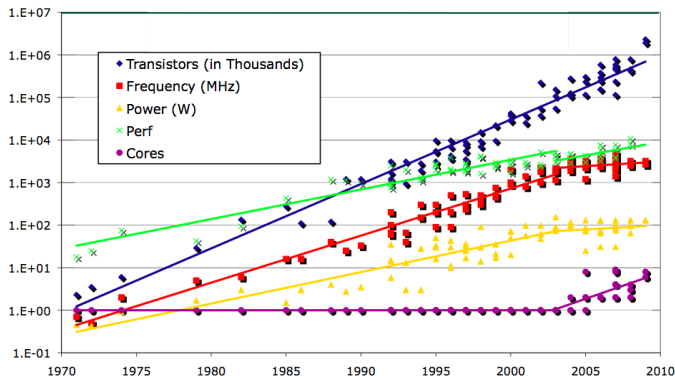
Non-Uniform Memory Architectures

Synchronization mechanisms

The memory consistency problem

Transistors, frequency, power, performance and ... cores!


An inflexion point in 2004 ... the power wall².




²Data from K. Olukotun, L. Hammond, H. Sutter, B. Smith, C. Batten, and K. Asanovic.

Multicore era

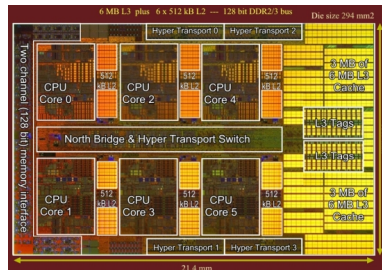
- ▶ The number of transistors on a chip is continuing to increase to accommodate multiple processors (cores) on a single chip
- ▶ Multicore = Chip Multi-Processor (CMP)



Multicore era (examples)




Intel's Nehalem i7 4-core (2 threads/core), 32/32 KB L1 and 256 KB L2 per core. Unified 4 or 8 MB L3




AMD's Istanbul 6-core, 512KB L2 per core and 6 MB shared L3

Multicore era (examples)




Sun's SPARC64-VIIIifx 8-core, 32/32K L1 per core and 5 MB shared L2. 2 GHz.. 760 M transistors



IBM's Power7 8-core (4 threads/core), 32/32K L1 and 256KB L2 per core. 32 MB shared L3. 2.4 GHz. 1.2 billion transistors

Multicore era (examples)

- ▶ Homogeneous multi-core systems include only identical cores, heterogeneous multi-core systems have cores with different functionalities, performances, ... (e.g. Intel Sandy Bridge)



Outline

Uniprocessor parallelism

Symmetric multi-processor architectures


Multicore architectures

Non-Uniform Memory Architectures

Synchronization mechanisms

The memory consistency problem

Non-Uniform Memory Architectures (NUMA)



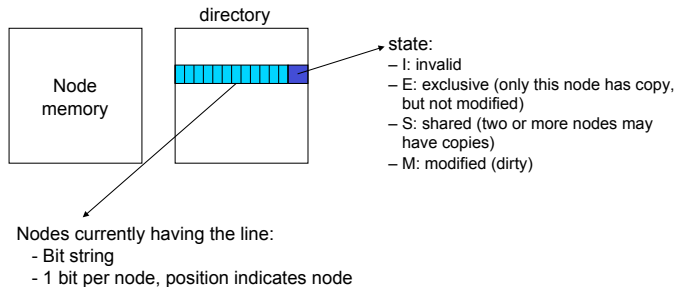
- Hub enables cache-coherent NUMA

Directory-based cache coherency protocol

- ▶ Scalability: coherence is maintained by point-to-point messages between the caches, not by broadcast mechanisms
 - ▶ **Home** node of a line: node (memory of the node) where the line is allocated (OS managed, for example first touch)
 - ▶ **Local** (requesting) node: node containing processor requesting a block
 - ▶ Other nodes may be involved: **Owner** node containing **dirty** copy or **Reader** nodes containing **clean** copies
- ▶ An additional structure is necessary to track the location of cache copies: **directory**


Directory-based cache coherency protocol

- ▶ Directory structure associated to each node: one directory entry per cache line of memory
 - ▶ State bits: they track the state of cache lines in its memory
 - ▶ Presence bits: if a block is shared, it must track which processors have it
 - ▶ If a block is valid only in one node, it should know which one




Directory-based cache coherency

- ▶ Example 1: read miss to clean line
 - ▶ Local node where the miss request originates
 - ▶ Home node where the cache line resides




Directory-based cache coherency

Example 2: Access to clean copy, with two readers, for write




Example 3: Access to dirty block for read




Front-Side Bus (FSB)

- ▶ Processor-northbridge (memory and I/O Hubs) connection (64-bit, 1-4 transfers/cycle, i.e. 6.4 GB/s @ 200 MHz)
- ▶ A number of processors can be connected to a single FSB sharing its total bandwidth
- ▶ Northbridge with several FSB ports ("snoop filters" to isolate FSB ports)



Quick Path Interconnect (QPI) and Hypertransport (HT)

- ▶ Point-to-point processor interconnects (16-bit data, 2 transfers/cycle, bidirectional, i.e. 25.6 GB/s @ 3.2 GHz)
- ▶ Connect several processors with integrated memory controllers and/or chipsets (IO): non-uniform memory access (NUMA)



Outline

Uniprocessor parallelism

Symmetric multi-processor architectures

Multicore architectures


Non-Uniform Memory Architectures

Synchronization mechanisms


The memory consistency problem

Shared memory: address space

Single process
sequential



Single process
multithreaded



Programmer needs

- ▶ Distribute work
- ▶ All threads can access data, heap and stacks
- ▶ Memory is not flat in a NUMA system
 - ▶ True and false sharing even more important
 - ▶ Data allocation and initialization sets the home node
 - ▶ Perform work according to data allocation to minimize data traffic
- ▶ Use synchronization mechanisms to avoid data races

Why synchronization?

- ▶ Needed to guarantee safety in the access to a shared-memory location or to signal a certain event
- ▶ Components:
 - ▶ Acquire method: acquire the right to do the synchronization
 - ▶ Waiting policy: busy wait, block/awake, wait for a while and then block, ...
 - ▶ Release method: allow others to proceed
- ▶ What's wrong with ...? (assume flag initialized to 0)

P1	P2
...	...
lock: ld r1, flag	lock: ld r1, flag
st flag, 1	st flag, 1
bnez r1, lock	bnez r1, lock
... // safe access	... // safe access
unlk: st flag, 0	unlk: st flag, 0
...	...

Support for synchronization at the architecture level

- ▶ Need hardware support to guarantee atomic (indivisible) instruction to fetch and update memory
 - ▶ User-level synchronization operations (e.g. locks, barriers, point-to-point, ...) using these primitives
- ▶ test-and-set: read value in location and set to 1

```
lock:  t&s r2, flag
      bnez r2, lock    // already locked?
```

- ▶ Atomic exchange: interchange of a value in a register with a value in memory
- ▶ fetch-and-op: read value in location and replace with increment/decrement

Support for synchronization at the architecture level

- ▶ Atomicity difficult or inefficient in large systems. Idea: assume optimistic atomic access
- ▶ Load-linked Store-conditional ll-sc
 - ▶ ll returns the current value of a memory location
 - ▶ sc stores a new value in that memory location if no updates have occurred to it since the ll; otherwise, the sc is guaranteed to fail

```
atomic: ll r1, location
        r2 = f (r1)           // f can be: "r2=1", "r2=r1", "r2=r1+1", ...
        sc location, r2
        beqz atomic           // Z flag set to 1 if atomicity has occurred
```

Other synchronization primitives

- ▶ How to implement a barrier synchronization primitive?
 - ▶ Processors block until all have reached it
 - ▶ Structure with fields {lock, counter, flag}

```
...
lock: t&s r2, barrier.lock      // acquire lock
     bnez r2, lock
     if (barr.counter == 0)
         barr.flag = 0          // reset flag if first
     mycount = barr.counter++;
     barr.lock = 0              // release lock
     if (mycount == P) {        // last to arrive?
         barr.counter = 0       // reset for next barrier
         barr.flag = 1          // release waiting processors
     } else
         while (barr.flag == 0) // busy wait for release
     ...
```

Reducing Synchronization Cost: test-test-and-set

- ▶ Acquire lock using test-and-set based synchronization

```
...  
lock: t&s r2, barrier.lock      // test and acquire lock if free  
      bnez r2, lock  
...
```

- ▶ test-test-and-set technique reduces the necessary memory bandwidth and coherence protocol operations required by a pure test-and-set based synchronization:
 - ▶ First, wait using a regular load instruction (lock will be cached)
 - ▶ Second, use test-and-set operation

```
...  
lock: ld r2, barrier.lock      // first test with regular load  
                                   // lock is cached meanwhile it is not updated  
      bnez r2, lock            // test if the lock is free  
      t&s r2, barrier.lock      // second test and acquire lock if STILL free  
      bnez r2, lock  
...
```

Reducing Synchronization Cost: test-test-and-set

- ▶ test-test-and-set **idea**³ can also help to reduce the synchronization cost of high level parallel programs
- ▶ Non optimized version : the synchronization is always done

```
acquire_lock(&lock);  
    if (value<CONSTANT)    // Test  
        value++;           // Set (Assign)  
release_lock(&lock);
```

- ▶ Optimized version : the synchronization is done if any chance of doing "Set" operation

```
if (value<CONSTANT)    // Test  
{  
    acquire_lock(&lock); // lock cost is only paid if necessary  
    if (value<CONSTANT) // Test  
        value++;       // Set (Assign)  
    release_lock(&lock);  
}
```

³Note that acquire_lock implementation may also use the test-test-and-set technique to reduce the synchronization cost of the operation

Outline

Uniprocessor parallelism

Symmetric multi-processor architectures

Multicore architectures

Non-Uniform Memory Architectures

Synchronization mechanisms

The memory consistency problem

Consistency

- ▶ In current systems, the compiler and hardware can freely reorder operations to different memory locations, as long as data/control dependences in sequential execution are guaranteed. This enables:
 - ▶ Compiler optimizations such as register allocation, code motion, loop transformations, ...
 - ▶ Hardware optimizations, such as pipelining, multiple issue, write buffer bypassing and forwarding, and lockup-free caches, ...

all of which lead to overlapping and reordering of memory operations

Consistency

- ▶ Will writes to different locations be seen in an order that makes sense, according to what is written in the source code?
- ▶ Example: two processors are synchronizing on a variable called flag. Assume A and flag are both initialized to 0

P1	P2
<pre>A=1; flag=1;</pre>	<pre>while (flag==0); /*spin*/ print A;</pre>

- ▶ What value does the programmer expect to be printed?

Memory consistency model

The memory consistency model ...

- ▶ Provides a formal specification of how the memory system will appear to the programmer ...
- ▶ By placing restrictions on the reordering of shared-memory operations

Sequential consistency, easy to understand but it may disallow many hardware and compiler optimizations that are possible in uniprocessors by enforcing a strict order among shared memory operations.

Memory consistency model

Relaxed consistency (weak), specifying regions of code within which shared-memory operations can be reordered

- ▶ fence machine instruction to force all pending memory operations to complete
- ▶ `#pragma omp flush` and other implicit points in OpenMP language

Different possibilities and implementations to be studied in *Multiprocessors* course

Parallelism (PAR)

Introduction to (shared-memory) parallel architectures

Eduard Ayguade, Josep Ramon Herrero and Daniel Jiménez
(`{eduard,josepr,djimenez}@ac.upc.edu`)

Computer Architecture Department
Universitat Politècnica de Catalunya

2014/15-Fall